

BASELINE SIMULATION

Max Welz

welz@ese.eur.nl

ECONOMETRIC INSTITUTE
ERASMUS SCHOOL OF ECONOMICS

August 25, 2021

We here describe a baseline simulation design. This simulation design is intended to be as simple as possible. The rationale is that if a certain estimator does not perform well in this simple framework, it will most likely also not perform well in more complicated frameworks. Concretely, this simulation design entails:

- No treatment effect heterogeneity;
- No noise;
- Randomized experiment;
- Perfect information;
- Linear data generating process;
- Identically and independently distributed Gaussian covariates;
- All deceasing individuals die at the exact same time.

Sampling of covariates. Let n be the number of observations and p be the number of covariates. The p -dimensional covariate vector of observation i , X_i , is standard normally distributed. The binary treatment assignment variable W_i assigns the treatment randomly. Mathematically, for all $i = 1, \dots, n$,

$$X_i \sim \mathcal{N}_p(0, I) \quad \text{and} \quad W_i \sim \text{Bernoulli}(0.5).$$

Predictive function. Suppose that $\beta_0 \in \mathbb{R}$ is an intercept term and that $\beta \in \mathbb{R}^p$ is a vector of coefficients associated with the covariates. Let $\theta \in \mathbb{R}$ be the *treatment effect parameter*, which governs the strength of the treatment effect. Since θ is a constant, there is no treatment effect heterogeneity. An effective treatment strategy W_i is characterized by $\theta \neq 0$. We assume that the *predictive function*, $\eta_i(W_i)$, a function of W_i , obeys

$$\eta_i(W_i) := \eta(W_i, X_i) = \beta_0 + X_i^\top \beta + \theta W_i.$$

Hence, the predictive function is assumed to be linear in its parameters (β_0, β, θ) .

Mortality risk. We emulate a medical trial. We are interested in the effect of treatment W_i on the (binary) mortality Y_i , where $Y_i = 1$ if individual i dies during the trial, and $Y_i = 0$ if it survives. The probability that individual i with treatment assignment status W_i dies is called the *mortality risk*, which is denoted by $P_i(W_i)$ and defined as

$$P_i(W_i) := P(W_i, X_i) = \mathbb{P}[Y_i = 1 | W_i, X_i] = F_{logistic}(\eta_i(W_i)),$$

where $F_{logistic}$ is the distribution function of the logistic distribution.

Potential Outcomes. The potential outcomes $Y_i(1)$ and $Y_i(0)$ denote the outcome if individual i is treated and if it is *not* treated, respectively. Since the outcome is binary, we assume that

$$Y_i(W_i) \sim \text{Bernoulli}(P_i(W_i)).$$

The *observed outcome* Y_i is defined by

$$Y_i = \begin{cases} Y_i(1) & \text{if } W_i = 1, \\ Y_i(0) & \text{if } W_i = 0. \end{cases}$$

Time at risk. We assume that the trial takes 10 years and that all deceasing individuals die after 7 years. That is, individual i 's *time-at-risk*, T_i , is given by

$$T_i = 7 + 3\mathbb{1}_{\{Y_i=0\}}.$$

Quantities of interest. We are interested in estimating the following quantities by using the observations $\{(X_i, Y_i, W_i, T_i)\}_{i=1}^n$:

- Treatment effect parameter θ ;
- 95% coverage of θ ;
- Average treatment effect: $\text{ATE} = \frac{1}{n} \sum_{i=1}^n (Y_i(1) - Y_i(0))$;
- Average relative treatment effect: $\text{ARTE} = \frac{\sum_{i=1}^n Y_i(0)}{\sum_{i=1}^n Y_i(0)}$;
- Calibration curve, both relative and absolute.

Parameter choices. We fix the parameters as follows.

- $n = 10,000$ and $p = 10$;
- $\theta = -2.5$ (that is, the treatment is effective);
- $\beta_0 = 0$, and, for $j = 1, \dots, p$,

$$\beta_j = \begin{cases} 1 & \text{if } j \text{ is even} \\ -1 & \text{if } j \text{ is odd.} \end{cases}$$

These choices lead to the “true” values $\text{ATE} \approx -0.25$ and $\text{ARTE} \approx 0.5$.